

Example of Object Detection with Detectron 2 based on the tutorial from:

https://colab.research.google.com/drive/16jcaJoc6bCFAQ96jDe2HwtXj7BMD_-m5#scrollTo=b2bjrfb2LDeo (https://colab.research.google.com/drive/16jcaJoc6bCFAQ96jDe2HwtXj7BMD_-m5#scrollTo=b2bjrfb2LDeo)

----- Environment Setup -----

This was ran on a Standard NV6_Promo (6 vcpus, 56 GiB memory) Instance on Azure which includes GPU's

Requirements:

Python 3.6

GPU - cuda-drivers (418.87.00-1)

Step 1 - The Cuda Driver on the Instance is not up-to-date, so an update is required to the NV6_Promo instance (as of 2019-01-01)

To update the Cuda Driver issue the following command:

```
sudo apt-get install cuda-drivers
```

A Reboot of the linux host is then required:

```
sudo reboot
```

Step 2 - Install the python libraries for Detectron 2

```
cd /data/anaconda/envs/py36/bin
```

```
./pip install torch==1.3.0+cu92 torchvision==0.4.1+cu92 -f https://download.pytorch.org/whl/torch\_stable.html (https://download.pytorch.org/whl/torch\_stable.html)
```

```
./pip install git+https://github.com/facebookresearch/fvcore.git (https://github.com/facebookresearch/fvcore.git)
```

```
git clone https://github.com/facebookresearch/detectron2 (https://github.com/facebookresearch/detectron2) detectron2_repo
```

```
./pip install -e detectron2_repo
```

```
./pip install 'git+https://github.com/facebookresearch/fvcore' (https://github.com/facebookresearch/fvcore)
```

```
./pip install opencv-python
```

```
./pip install "git+https://github.com/philferriere/cocoapi.git#egg=pycocotools&subdirectory=PythonAPI" (https://github.com/philferriere/cocoapi.git#egg=pycocotools&subdirectory=PythonAPI)"
```

If using Jupyter a restart of the kernel is necessary for the notebook to the new packages

Create Cell Dataset which will be used to for object detection, also create a project specific instance of the detectron2_repo by running the output of the next cell in the linux command line

```
In [1]: import os
dirpath = os.getcwd()
print("cd %s" % dirpath)
print("git clone https://github.com/facebookresearch/detectron2 detectron2_repo")

cd /data/home/ben_dk_ds/notebooks/hw
git clone https://github.com/facebookresearch/detectron2 detectron2_repo
```

Need to unzip files on first pass

```
In [2]: #import zipfile
#with zipfile.ZipFile("cell_obj_detection_images.zip", 'r') as zip_ref:
#    zip_ref.extractall("./")
```

```
In [3]: # You may need to restart your runtime prior to this, to let your installation take effect
# Some basic setup
# Setup detectron2 logger
import detectron2
from detectron2.utils.logger import setup_logger
setup_logger()

# import some common libraries
import numpy as np
import cv2
#from google.colab.patches import cv2_imshow

# import some common detectron2 utilities
from detectron2.engine import DefaultPredictor
from detectron2.config import get_cfg
from detectron2.utils.visualizer import Visualizer
from detectron2.data import MetadataCatalog

import os
import numpy as np
import json
from detectron2.structures import BoxMode
import itertools
%matplotlib inline
```

```
In [4]: # write a function that loads the dataset into detectron2's standard format
def get_cell_dicts(img_dir):
    json_file = os.path.join(img_dir, "via_region_data.json")
    with open(json_file) as f:
        imgs_anns = json.load(f)

    dataset_dicts = []
    for _, v in imgs_anns.items():
        record = {}

        filename = os.path.join(img_dir, v["filename"])
        height, width = cv2.imread(filename).shape[:2]

        record["file_name"] = filename
        record["height"] = height
        record["width"] = width

        annos = v["regions"]
        objs = []
        for _, anno in annos.items():
            assert not anno["region_attributes"]
            anno = anno["shape_attributes"]
            px = anno["all_points_x"]
            py = anno["all_points_y"]
            poly = [(x + 0.5, y + 0.5) for x, y in zip(px, py)]
            poly = list(itertools.chain.from_iterable(poly))

            obj = {
                "bbox": [np.min(px), np.min(py), np.max(px), np.max(py)],
                "bbox_mode": BoxMode.XYXY_ABS,
                "segmentation": [poly],
                "category_id": 0,
                "iscrowd": 0
            }
            objs.append(obj)
        record["annotations"] = objs
        dataset_dicts.append(record)
    return dataset_dicts
```



```
In [5]: source_dir = "upload/"
train_dir = "%strain" % source_dir
valid_dir = "%svalid" % source_dir
from detectron2.data import DatasetCatalog, MetadataCatalog
for d in ["train", "valid"]:
    DatasetCatalog.register(source_dir + d, lambda d=d: get_cell_dicts(source_dir + d))
    MetadataCatalog.get(source_dir + d).set(thing_classes=["cell"])
cell_metadata = MetadataCatalog.get(train_dir)
```

```
In [6]: dataset_dicts = get_cell_dicts(train_dir)
```



```
In [7]: %pylab inline
import random
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

dataset_dicts = get_cell_dicts(valid_dir)
for d in random.sample(dataset_dicts, 3):
    print(d["file_name"])
    img = cv2.imread(d["file_name"])
    imgplot = plt.imshow(img)
    plt.show()
    visualizer = Visualizer(img[:, :, ::-1], metadata=cell_metadata, scale=1)
    vis = visualizer.draw_dataset_dict(d)
    imgplot = plt.imshow(vis.get_image()[:, :, ::-1])
    plt.show()
```



Populating the interactive namespace from numpy and matplotlib
upload/valid/p232_0_A06F00d2.jpg



upload/valid/p388_0_A04F00d2.jpg



upload/valid/p544_0_A04F00d2.jpg



```
In [20]: from detectron2.engine import DefaultTrainer
        from detectron2.config import get_cfg

        cfg = get_cfg()
        cfg.merge_from_file("./detectron2_repo/configs/COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml")
        cfg.DATASETS.TRAIN = (train_dir,)
        cfg.DATASETS.TEST = () # no metrics implemented for this dataset
        cfg.DATALOADER.NUM_WORKERS = 2
        cfg.MODEL.WEIGHTS = "detectron2://COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x/137849600/model_final_f10
        217.pkl" # initialize from model zoo
        cfg.SOLVER.IMS_PER_BATCH = 3
        cfg.SOLVER.BASE_LR = 0.00025
        cfg.SOLVER.MAX_ITER = 500 # 300 iterations seems good enough, but you can certainly train longer
        cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 256 # faster, and good enough for this toy dataset
        cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1 # only has one class (cell)
        cfg.MODEL.MASK_ON = True

        os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
        trainer = DefaultTrainer(cfg)
        trainer.resume_or_load(resume=False)
        trainer.train()
```

WARNING [11/24 21:25:13 d2.config.compat]: Config './detectron2_repo/configs/COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml' has no VERSION. Assuming it to be compatible with latest v2.

[11/24 21:25:14 d2.engine.defaults]: Model:

```
GeneralizedRCNN(
  (backbone): FPN(
    (fpn_lateral2): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (fpn_lateral3): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (fpn_lateral4): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (fpn_lateral5): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
    (fpn_output5): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (top_block): LastLevelMaxPool()
  )
  (bottom_up): ResNet(
    (stem): BasicStem(
      (conv1): Conv2d(
        3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
    )
    (res2): Sequential(
      (0): BottleneckBlock(
        (shortcut): Conv2d(
          64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
        )
        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
      )
      (conv1): Conv2d(
        64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      (conv3): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (1): BottleneckBlock(
      (conv1): Conv2d(
        256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      (conv3): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
    (2): BottleneckBlock(
      (conv1): Conv2d(
        256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      (conv2): Conv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
      (conv3): Conv2d(
        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
    )
  )
  (res3): Sequential(
    (0): BottleneckBlock(
      (shortcut): Conv2d(
        256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False
      )
      (norm): FrozenBatchNorm2d(num_features=512, eps=1e-05)
    )
    (conv1): Conv2d(
      256, 128, kernel_size=(1, 1), stride=(2, 2), bias=False
    )
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
    (conv2): Conv2d(
      128, 128, kernel_size=(1, 1), stride=(1, 1), bias=False
    )
    (norm): FrozenBatchNorm2d(num_features=128, eps=1e-05)
  )
)
```

'roi_heads.box_predictor.cls_score.weight' has shape (81, 1024) in the checkpoint but (2, 1024) in the model! Skipped.
'roi_heads.box_predictor.cls_score.bias' has shape (81,) in the checkpoint but (2,) in the model! Skipped.
'roi_heads.box_predictor.bbox_pred.weight' has shape (320, 1024) in the checkpoint but (4, 1024) in the model! Skipped.
'roi_heads.box_predictor.bbox_pred.bias' has shape (320,) in the checkpoint but (4,) in the model! Skipped.
'roi_heads.mask_head.predictor.weight' has shape (80, 256, 1, 1) in the checkpoint but (1, 256, 1, 1) in the model! Skipped.
'roi_heads.mask_head.predictor.bias' has shape (80,) in the checkpoint but (1,) in the model! Skipped.


```

[11/24 21:25:22 d2.engine.train_loop]: Starting training from iteration 0
[11/24 21:25:45 d2.utils.events]: eta: 0:09:24 iter: 19 total_loss: 5.742 loss_cls: 0.659 loss_box_re
g: 0.634 loss_mask: 0.690 loss_rpn_cls: 3.072 loss_rpn_loc: 0.620 time: 1.1619 data_time: 0.0365 l
r: 0.000005 max_mem: 3792M
[11/24 21:26:08 d2.utils.events]: eta: 0:08:56 iter: 39 total_loss: 3.520 loss_cls: 0.614 loss_box_re
g: 0.675 loss_mask: 0.683 loss_rpn_cls: 0.935 loss_rpn_loc: 0.590 time: 1.1527 data_time: 0.0311 l
r: 0.000010 max_mem: 3792M
[11/24 21:26:30 d2.utils.events]: eta: 0:08:24 iter: 59 total_loss: 2.659 loss_cls: 0.575 loss_box_re
g: 0.689 loss_mask: 0.670 loss_rpn_cls: 0.150 loss_rpn_loc: 0.535 time: 1.1381 data_time: 0.0300 l
r: 0.000015 max_mem: 3792M
[11/24 21:26:52 d2.utils.events]: eta: 0:07:55 iter: 79 total_loss: 2.398 loss_cls: 0.519 loss_box_re
g: 0.694 loss_mask: 0.654 loss_rpn_cls: 0.111 loss_rpn_loc: 0.419 time: 1.1262 data_time: 0.0301 l
r: 0.000020 max_mem: 3792M
[11/24 21:27:15 d2.utils.events]: eta: 0:07:33 iter: 99 total_loss: 2.429 loss_cls: 0.467 loss_box_re
g: 0.680 loss_mask: 0.636 loss_rpn_cls: 0.117 loss_rpn_loc: 0.501 time: 1.1270 data_time: 0.0309 l
r: 0.000025 max_mem: 3792M
[11/24 21:27:38 d2.utils.events]: eta: 0:07:12 iter: 119 total_loss: 2.233 loss_cls: 0.441 loss_box_r
eg: 0.660 loss_mask: 0.618 loss_rpn_cls: 0.107 loss_rpn_loc: 0.407 time: 1.1300 data_time: 0.0316 l
r: 0.000030 max_mem: 3792M
[11/24 21:28:01 d2.utils.events]: eta: 0:06:51 iter: 139 total_loss: 2.254 loss_cls: 0.420 loss_box_r
eg: 0.663 loss_mask: 0.603 loss_rpn_cls: 0.102 loss_rpn_loc: 0.450 time: 1.1342 data_time: 0.0306 l
r: 0.000035 max_mem: 3792M
[11/24 21:28:23 d2.utils.events]: eta: 0:06:27 iter: 159 total_loss: 2.194 loss_cls: 0.394 loss_box_r
eg: 0.640 loss_mask: 0.592 loss_rpn_cls: 0.088 loss_rpn_loc: 0.489 time: 1.1319 data_time: 0.0309 l
r: 0.000040 max_mem: 3792M
[11/24 21:28:46 d2.utils.events]: eta: 0:06:04 iter: 179 total_loss: 2.099 loss_cls: 0.367 loss_box_r
eg: 0.648 loss_mask: 0.574 loss_rpn_cls: 0.089 loss_rpn_loc: 0.406 time: 1.1326 data_time: 0.0316 l
r: 0.000045 max_mem: 3792M
[11/24 21:29:09 d2.utils.events]: eta: 0:05:41 iter: 199 total_loss: 2.122 loss_cls: 0.358 loss_box_r
eg: 0.664 loss_mask: 0.562 loss_rpn_cls: 0.078 loss_rpn_loc: 0.441 time: 1.1326 data_time: 0.0303 l
r: 0.000050 max_mem: 3792M
[11/24 21:29:31 d2.utils.events]: eta: 0:05:18 iter: 219 total_loss: 1.974 loss_cls: 0.340 loss_box_r
eg: 0.633 loss_mask: 0.551 loss_rpn_cls: 0.086 loss_rpn_loc: 0.373 time: 1.1322 data_time: 0.0313 l
r: 0.000055 max_mem: 3792M
[11/24 21:29:54 d2.utils.events]: eta: 0:04:56 iter: 239 total_loss: 1.994 loss_cls: 0.322 loss_box_r
eg: 0.591 loss_mask: 0.529 loss_rpn_cls: 0.083 loss_rpn_loc: 0.396 time: 1.1350 data_time: 0.0333 l
r: 0.000060 max_mem: 3792M
[11/24 21:30:17 d2.utils.events]: eta: 0:04:33 iter: 259 total_loss: 1.905 loss_cls: 0.318 loss_box_r
eg: 0.600 loss_mask: 0.512 loss_rpn_cls: 0.073 loss_rpn_loc: 0.403 time: 1.1346 data_time: 0.0305 l
r: 0.000065 max_mem: 3792M
[11/24 21:30:40 d2.utils.events]: eta: 0:04:10 iter: 279 total_loss: 1.840 loss_cls: 0.308 loss_box_r
eg: 0.597 loss_mask: 0.487 loss_rpn_cls: 0.072 loss_rpn_loc: 0.394 time: 1.1336 data_time: 0.0308 l
r: 0.000070 max_mem: 3792M
[11/24 21:31:02 d2.utils.events]: eta: 0:03:47 iter: 299 total_loss: 1.735 loss_cls: 0.286 loss_box_r
eg: 0.559 loss_mask: 0.471 loss_rpn_cls: 0.073 loss_rpn_loc: 0.349 time: 1.1329 data_time: 0.0306 l
r: 0.000075 max_mem: 3792M
[11/24 21:31:25 d2.utils.events]: eta: 0:03:25 iter: 319 total_loss: 1.751 loss_cls: 0.292 loss_box_r
eg: 0.541 loss_mask: 0.449 loss_rpn_cls: 0.069 loss_rpn_loc: 0.394 time: 1.1344 data_time: 0.0308 l
r: 0.000080 max_mem: 3792M
[11/24 21:31:48 d2.utils.events]: eta: 0:03:03 iter: 339 total_loss: 1.707 loss_cls: 0.284 loss_box_r
eg: 0.507 loss_mask: 0.423 loss_rpn_cls: 0.080 loss_rpn_loc: 0.434 time: 1.1360 data_time: 0.0328 l
r: 0.000085 max_mem: 3792M
[11/24 21:32:11 d2.utils.events]: eta: 0:02:40 iter: 359 total_loss: 1.571 loss_cls: 0.262 loss_box_r
eg: 0.477 loss_mask: 0.414 loss_rpn_cls: 0.057 loss_rpn_loc: 0.323 time: 1.1348 data_time: 0.0304 l
r: 0.000090 max_mem: 3792M
[11/24 21:32:34 d2.utils.events]: eta: 0:02:17 iter: 379 total_loss: 1.553 loss_cls: 0.273 loss_box_r
eg: 0.447 loss_mask: 0.390 loss_rpn_cls: 0.068 loss_rpn_loc: 0.372 time: 1.1364 data_time: 0.0315 l
r: 0.000095 max_mem: 3792M
[11/24 21:32:57 d2.utils.events]: eta: 0:01:55 iter: 399 total_loss: 1.435 loss_cls: 0.252 loss_box_r
eg: 0.398 loss_mask: 0.380 loss_rpn_cls: 0.055 loss_rpn_loc: 0.327 time: 1.1376 data_time: 0.0321 l
r: 0.000100 max_mem: 3792M
[11/24 21:33:20 d2.utils.events]: eta: 0:01:32 iter: 419 total_loss: 1.435 loss_cls: 0.253 loss_box_r
eg: 0.408 loss_mask: 0.369 loss_rpn_cls: 0.070 loss_rpn_loc: 0.318 time: 1.1374 data_time: 0.0301 l
r: 0.000105 max_mem: 3792M
[11/24 21:33:43 d2.utils.events]: eta: 0:01:09 iter: 439 total_loss: 1.312 loss_cls: 0.232 loss_box_r
eg: 0.348 loss_mask: 0.373 loss_rpn_cls: 0.060 loss_rpn_loc: 0.294 time: 1.1372 data_time: 0.0313 l
r: 0.000110 max_mem: 3792M
[11/24 21:34:05 d2.utils.events]: eta: 0:00:46 iter: 459 total_loss: 1.389 loss_cls: 0.250 loss_box_r
eg: 0.370 loss_mask: 0.363 loss_rpn_cls: 0.056 loss_rpn_loc: 0.344 time: 1.1368 data_time: 0.0305 l
r: 0.000115 max_mem: 3792M
[11/24 21:34:28 d2.utils.events]: eta: 0:00:23 iter: 479 total_loss: 1.335 loss_cls: 0.229 loss_box_r
eg: 0.341 loss_mask: 0.361 loss_rpn_cls: 0.064 loss_rpn_loc: 0.335 time: 1.1377 data_time: 0.0313 l
r: 0.000120 max_mem: 3792M
[11/24 21:34:52 d2.utils.events]: eta: 0:00:01 iter: 499 total_loss: 1.305 loss_cls: 0.221 loss_box_r
eg: 0.336 loss_mask: 0.357 loss_rpn_cls: 0.046 loss_rpn_loc: 0.295 time: 1.1374 data_time: 0.0314 l
r: 0.000125 max_mem: 3792M
[11/24 21:34:55 d2.engine.hooks]: Overall training speed: 497 iterations in 0:09:26 (1.1397 s / it)
[11/24 21:34:55 d2.engine.hooks]: Total training time: 0:09:31 (0:00:04 on hooks)

```

Out[20]: OrderedDict()

```
In [21]: cfg.MODEL.WEIGHTS = os.path.join(cfg.OUTPUT_DIR, "model_final.pth")
         cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.7    # set the testing threshold for this model
         cfg.DATASETS.TEST = (valid_dir, )
         predictor = DefaultPredictor(cfg)
```


```

In [22]: %pylab inline
import random
import matplotlib.pyplot as plt
import matplotlib.image as mpimg


from detectron2.utils.visualizer import ColorMode
dataset_dicts = get_cell_dicts(valid_dir)
for d in random.sample(dataset_dicts, 3):
    im = cv2.imread(d["file_name"])
    fig = plt.figure(figsize=(10,10))
    imgplot = plt.imshow(im)
    plt.show()
    outputs = predictor(im)
    v = Visualizer(im[:, :, :-1],
                   metadata=cell_metadata,
                   scale=3,
                   instance_mode=ColorMode.IMAGE_BW    # remove the colors of unsegmented pixels
    )
    #print(outputs["instances"].get_fields()) #['pred_masks']
    v = v.draw_instance_predictions(outputs["instances"].to("cpu"))
    print(d["file_name"])
    fig = plt.figure(figsize=(10,10))
    imgplot = plt.imshow(v.get_image()[:, :, :-1])
    plt.show()


```

Populating the interactive namespace from numpy and matplotlib





upload/valid/p180_0_A04F00d2.jpg





upload/valid/p310_0_A01F00d2.jpg





upload/valid/p492_0_A04F00d2.jpg

